A Decision-Making Framework for CI Failure Diagnosis with Incomplete Evidence and Uncertain Root Causes

Pradnyawant M. Gote
AI Engineer @ StatGlow
pradnyawantgote100@gmail.com

Abstract

Continuous integration (CI) systems reveal where a workflow stopped, but the failed stage alone rarely identifies why it failed. This paper asks whether a diagnostic agent can decide which evidence to collect next instead of making a one-shot root-cause prediction. We formulate CI diagnosis as a finite-horizon Bayesian decision process over four operational categories: Code, Test, Dependency, and CI/Config. A stage-conditioned prior is initialized from 309 mapped records in a 375-failure GitHub Actions dataset. The agent then selects read-only diagnostic actions using expected Bayes-risk reduction per unit cost, updates its belief from binary outcomes, and either reports a category or refers the case for human review. In a deterministic simulation of 40 balanced cases, the proposed policy produced 17 correct and one incorrect automatic report, corresponding to 45.0% coverage and 94.44% selective accuracy. Its configured total decision cost was 589.0, compared with 643.0 for information-gain-per-cost, 603.5 for the lowest-cost policy, and 1020.0 for a stage-MAP baseline. The stage baseline reported every case but made 22 errors. However, the proposed policy referred all Test cases and eight of ten Code cases, and its single error was reported with 0.9334 confidence. The study therefore demonstrates an inspectable cost–risk trade-off inside a controlled simulator; it does not establish production accuracy because action likelihoods, outcomes, and costs are modeled rather than measured.

1 Introduction

Continuous integration provides rapid feedback on whether software changes can be built, tested, and delivered, but a failed workflow is an observation rather than an explanation [Shahin *et al.*, 2017]. Source-code defects, failing tests, dependency changes, and workflow or environment faults can interrupt the same pipeline stage. In a recent empirical study, Zheng et al. [Zheng *et al.*, 2026] manually examined 375 failed GitHub Actions runs from 260 Java projects and reported a broad taxonomy of causes. Other evidence shows that some build failures are unrelated to the triggering patch [Huang *et al.*, 2026]. These findings make one-shot diagnosis from the stage or latest commit a weak foundation for action.

Automated systems increasingly extract useful log evidence and generate failure explanations. LogSage combines log processing, LLM-based diagnosis, retrieval, and remediation [Xu *et al.*, 2025], while Auto-Diagnose presents log-grounded integration-test summaries inside Google's review workflow [Ziftci *et al.*, 2026]. These systems address how evidence can be interpreted. The complementary question studied here is how an agent should choose the *next* investigation when several causes remain plausible and diagnostic checks have different costs.

We model that choice as sequential, decision-theoretic troubleshooting [Heckerman *et al.*, 1995]. The agent maintains a belief over Code, Test, Dependency, and CI/Config, predicts the possible outcomes of each available read-only test, and greedily selects the action with the largest expected reduction in reporting risk per unit cost. The belief is updated after each outcome, and the episode ends with an automatic report only when a confidence and loss gate is met; otherwise the evidence is referred to a person. This design follows Bayesian experimental-design principles [Rainforth *et al.*, 2024], but it is a one-step policy rather than an optimal multi-step planner.

The contribution is threefold. First, the study gives a transparent probability-and-decision model for evidence acquisition, reporting, and referral. Second, it evaluates four policies on the same balanced cases and pre-generated outcome bank using coverage, conditional accuracy, probability quality, diagnostic effort, and asymmetric decision cost. Third, it audits the belief trajectories behind the observed high-confidence errors. The scope is deliberately limited: historical records provide labels and stage context, whereas diagnostic likelihoods, binary outcomes, and costs are declared or simulated. The reported numbers describe this controlled model, not unseen repositories or live incident response.

2 Related Work

Empirical studies show that a failed run can have several explanations. Zheng et al. [Zheng *et al.*, 2026] provide the GitHub Actions taxonomy and replication data used in this

study. Huang et al. [Huang *et al.*, 2026] examine failures unrelated to the current patch, while Huang and Lin [Huang and Lin, 2026] study reruns whose outcomes can change without repository modifications. Together, these studies support separating the observed failure, the triggering change, and the underlying cause.

Prior tools operate at finer granularity than the four-category belief used here. BuildSheriff groups related CI test failures using change and failure information [Zhang *et al.*, 2022]; UniLoc retrieves suspicious source and build-script files from failure logs [Hassan *et al.*, 2023]; and LogChunks provides 797 annotated Travis CI logs for evaluating failure-evidence retrieval [Brandt *et al.*, 2020]. These methods could become evidence producers for a future agent, whereas the present experiment assumes binary outcomes and studies their selection order.

Recent work uses LLMs to summarize or resolve CI failures. Developer evaluation of GitHub Actions explanations found encouraging results for simpler logs but continued difficulty with complex scenarios [Valenzuela-Toledo *et al.*, 2025]. LogSage evaluates an end-to-end diagnosis and remediation pipeline [Xu *et al.*, 2025], and Auto-Diagnose reports deployment evidence for integration-test assistance at Google [Ziftci *et al.*, 2026]. The current agent is narrower: it neither parses raw logs nor changes code or runtime state. It chooses among authorized investigations and records why a report or referral was made.

Decision-theoretic troubleshooting combines diagnostic probabilities with the costs of observations and wrong decisions [Heckerman *et al.*, 1995]. Belief states similarly summarize uncertainty in partially observable decision processes [Kaelbling *et al.*, 1998]. Lindley's expected information measure [Lindley, 1956] and modern Bayesian experimental design [Rainforth *et al.*, 2024] motivate choosing observations before their outcomes are known. Our model applies these established ideas to a CI simulation, but differs from general POMDP planning by using a greedy one-step score and a fixed state space. Its novelty lies in the inspectable connection among evidence value, asymmetric reporting loss, investigation cost, and human referral, not in a new Bayesian inference rule.

3 Agent Design

The agent treats continuous integration (CI) failure diagnosis as a sequence of decisions under uncertainty. It maintains competing explanations for a failure, chooses a diagnostic action, and revises its belief after observing the result. Investigation ends with either a root-cause report or referral to a human reviewer. This formulation draws on decision-theoretic troubleshooting and belief-state reasoning under partial observability [Heckerman *et al.*, 1995; Kaelbling *et al.*, 1998]. The proposed policy uses a one-step estimate of the reduction in reporting loss to choose the next action as illustrated in Figure 1.

3.1 State representation and evidence intake

The hidden state $s^\star$ denotes the failure category for the current incident. The model uses

$$\mathcal{S} = \{\text{Code}, \text{Test}, \text{Dependency}, \text{CI/Config}\}. \tag{1}$$

These are operational categories inherited from the project's mapping of historical failure labels. In particular, a failure assigned to *Test* need not originate in the test code itself. Each incident is assigned one category, and diagnostic actions are assumed to leave that category unchanged during an investigation. Simultaneous faults are outside this representation.

The evidence-intake interface organizes the failed stage, logs, test output, recent code or dependency changes, and CI environment context. In the current notebook, the agent receives a case identifier and normalized stage initially, followed by binary diagnostic outcomes. Parsing the raw artifacts is therefore an interface requirement rather than an implemented log-analysis capability. The identifier is used for tracking and does not affect inference.

The implemented normalizer uses seven labels: Build, Test, Setup / Dependency, Quality / Analysis, Deploy / Publish, Ambiguous, and Unknown. These refine the build–test–deploy overview in the architecture. Missing stage information maps to Unknown, while unmatched descriptions map to Ambiguous. Let z be the normalized stage and $\mathcal{Z}$ this seven-label vocabulary.

3.2 Historical initialization

The source workbook comes from the replication package of Zheng et al. [Zheng *et al.*, 2026]. The project's mapping retains 309 of its 375 records and excludes 66 records outside the four modeled categories. The retained counts for Code, Test, Dependency, and CI/Config are 54, 120, 91, and 44, respectively. These counts describe the selected source population rather than the prevalence of failures across all CI systems.

Let n_s denote the number of retained records in category s, and $n_{s,z}$ the number with category s and stage z. Laplace smoothing gives

$$\widehat{P}(s) = \frac{n_s + \alpha}{N + \alpha|\mathcal{S}|}, \tag{2}$$

$$\widehat{P}(z \mid s) = \frac{n_{s,z} + \beta}{n_s + \beta|\mathcal{Z}|}, \tag{3}$$

where $N = 309$ and $\alpha = \beta = 1$. The initial belief incorporates the observed stage once:

$$b_0(s) = \frac{\widehat{P}(z \mid s)\widehat{P}(s)}{\sum_{u \in \mathcal{S}} \widehat{P}(z \mid u)\widehat{P}(u)}. \tag{4}$$

Smoothing keeps a category or stage combination from receiving zero probability solely because it is absent from the finite source sample.

3.3 Diagnostic actions and belief updates

The action catalogue covers log inspection, failed-stage inspection, comparison with the last successful run, recent changes, dependencies, CI configuration, previous incidents, Docker images or caches, hosting, and server or session state.

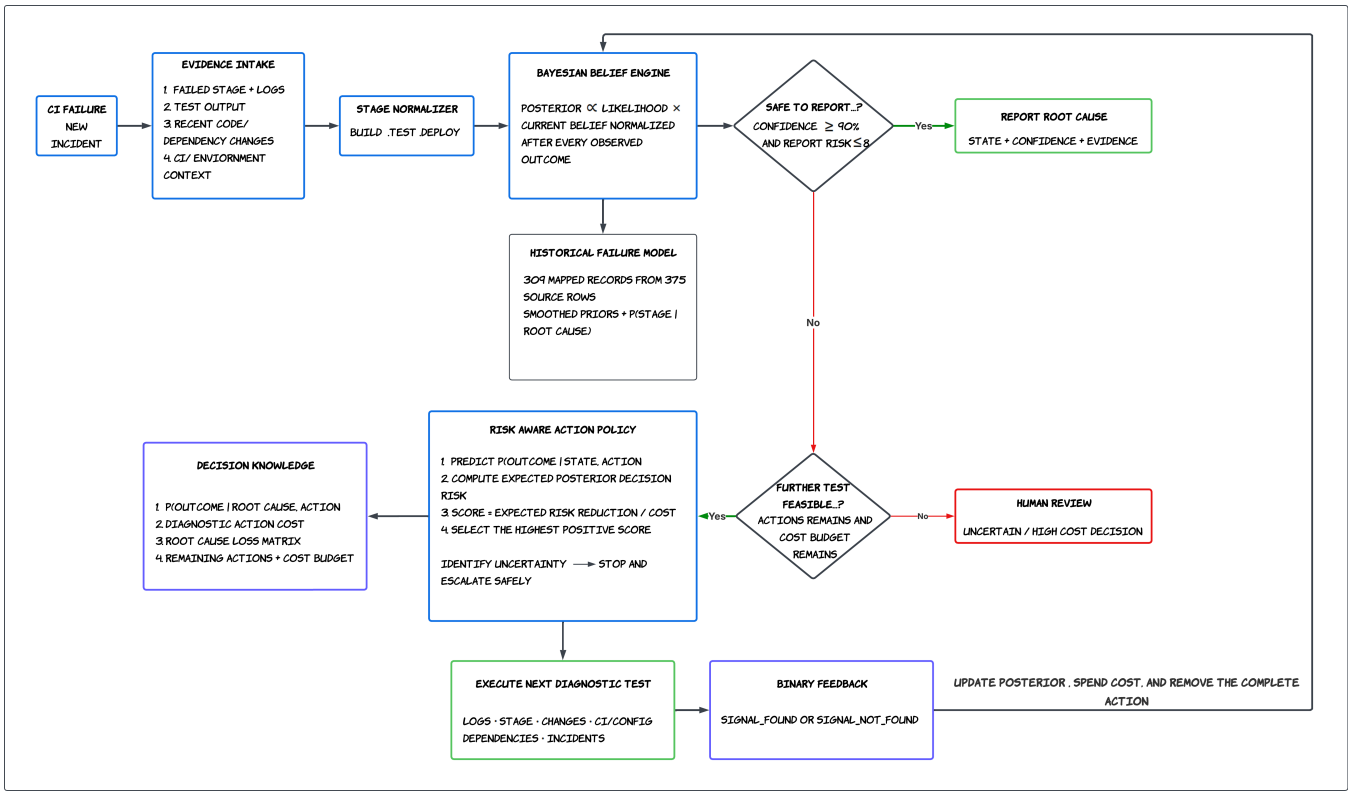


Figure 1: Conceptual architecture of agentic probability behavior and decision making.

Each action a has a positive cost $c(a)$ and a set of eligible stages. Costs range from 1 to 2.5 relative effort units; for example, dependency inspection costs 1 and CI configuration inspection costs 2.5. These are declared modeling values rather than measured durations or monetary costs.

An action returns $o \in \{0, 1\}$, where 1 denotes `SIGNAL_FOUND` and 0 denotes `SIGNAL_NOT_FOUND`. Its observation model is

$$\ell_a(o \mid s) = \theta_{s,a}^{o}(1 - \theta_{s,a})^{1-o}, \qquad \theta_{s,a} = P(o = 1 \mid s, a). \tag{5}$$

The diagnostic likelihoods $\theta_{s,a}$ are specified assumptions; the historical workbook does not provide the action outcomes needed to estimate them. A positive signal is evidence whose meaning depends on the action, not a direct declaration of the true category.

At step t, the belief $b_t(s)$ represents the posterior probability of category s given the stage and outcomes observed so far. For a candidate action, the predictive outcome probability and hypothetical posterior are

$$q_t(o \mid a) = \sum_{s \in \mathcal{S}} b_t(s)\ell_a(o \mid s), \tag{6}$$

$$b_t^{a,o}(s) = \frac{b_t(s)\ell_a(o \mid s)}{q_t(o \mid a)}. \tag{7}$$

The posterior is defined for outcomes with $q_t(o \mid a) > 0$; zero-probability branches contribute nothing to expectations. After observing o_t from the selected action a_t, the agent sets $b_{t+1} = b_t^{a_t,o_t}$.

Using fixed action likelihoods assumes that outcomes are conditionally independent given the hidden category and chosen actions. An additional stage inspection must therefore provide a new clue, rather than count the already observed stage a second time. Dependencies between signals from overlapping logs or changes are a limitation of this model.

3.4 Reporting loss and acceptance criteria

Let $L(s, d)$ be the terminal loss of reporting category d when the true category is s. The configured loss matrix is

$$L = \begin{pmatrix} 0 & 40 & 40 & 40 \\ 25 & 0 & 25 & 25 \\ 45 & 45 & 0 & 45 \\ 65 & 65 & 65 & 0 \end{pmatrix}. \tag{8}$$

Rows denote true categories and columns reported categories, both ordered as Code, Test, Dependency, and CI/Config. A wrong report incurs the true-category miss cost plus a fixed add-on of 5. The values express the project's assumed consequences of misdiagnosis; they are not empirical estimates of incident impact.

The expected loss of a report and the minimum reporting risk are

$$r(d \mid b) = \sum_{s \in \mathcal{S}} b(s) L(s, d), \tag{9}$$

$$d_t^{\text{risk}} \in \arg\min_{d \in \mathcal{S}} r(d \mid b_t), \qquad R(b_t) = r(d_t^{\text{risk}} \mid b_t). \tag{10}$$

Here R concerns root-cause reports only; it does not include human review as a competing report. Define the maximum a posteriori (MAP) category as $d_t^{\text{MAP}} \in \arg\max_s b_t(s)$. Automatic reporting requires

$$\begin{aligned} d_t^{\text{MAP}} &= d_t^{\text{risk}}, \\ b_t(d_t^{\text{MAP}}) &\geq 0.90, \\ R(b_t) &\leq 8. \end{aligned} \tag{11}$$

The risk threshold is expressed in the declared loss units, rather than as a percentage. With this particular matrix, 90% confidence already implies agreement between the MAP and minimum-risk categories and a reporting loss of at most $65(1 - 0.90) = 6.5$. The other two checks are consequently redundant at the current settings, although they remain explicit configuration checks. The risk model still affects action selection at lower confidence.

3.5 Selecting the next diagnostic action

Let $\mathcal{U}_t$ be the unused actions and B_t the remaining diagnostic budget. After t executed actions, the feasible set is

$$\mathcal{F}_t = \{a \in \mathcal{U}_t : a \text{ is eligible for } z,\ c(a) \leq B_t,\ t < 8\}, \tag{12}$$

with $B_0 = 11$. If reporting is not yet permitted, the agent evaluates how each feasible action could change the best subsequent report:

$$\overline{R}_t(a) = \sum_{o \in \{0,1\}} q_t(o \mid a) R(b_t^{a,o}), \tag{13}$$

$$\Delta R_t(a) = R(b_t) - \overline{R}_t(a), \tag{14}$$

$$V_t(a) = \frac{\Delta R_t(a)}{c(a)}. \tag{15}$$

The minimum-risk report is recomputed separately for each hypothetical outcome. This allows an action to be valuable because it could change the reporting decision, even if it does not produce the largest reduction in overall uncertainty.

Following the architecture, the design selects

$$a_t \in \arg\max_{a \in \mathcal{F}_t} V_t(a) \quad \text{only if} \quad \max_{a \in \mathcal{F}_t} V_t(a) > 0. \tag{16}$$

Ties are resolved by larger expected risk reduction, then lower cost, then action name. The positive-score condition is an explicit design requirement: the current notebook ranks candidates but does not yet reject an action whose score is zero. If the feasible set is empty or no positive score remains, the design refers the incident for human review.

This is a greedy, one-step policy. It does not optimize complete sequences of tests or guarantee the minimum total investigation cost. In particular, a positive risk reduction does not by itself establish that the improvement outweighs the action's cost under a common utility scale.

For comparison, the EIG-per-cost policy uses expected information gain:

$$I_t(a) = H(b_t) - \sum_{o \in \{0,1\}} q_t(o \mid a) H(b_t^{a,o}), \tag{17}$$

where $H(b) = -\sum_s b(s) \log_2 b(s)$, with $0 \log_2 0 = 0$. That policy ranks actions by $I_t(a)/c(a)$; the proposed policy ranks them by Equation (15), which accounts for the specified consequences of reporting errors.

3.6 Execution, escalation, and audit trail

After executing a_t and updating the posterior, the agent updates its resources:

$$B_{t+1} = B_t - c(a_t), \qquad \mathcal{U}_{t+1} = \mathcal{U}_t \setminus \{a_t\}. \tag{18}$$

It then checks the reporting conditions again, including after the last permitted action. An accepted report contains the selected category, posterior confidence, and supporting diagnostic outcomes. Otherwise, exhaustion of the action count, affordable eligible actions, or positive scores leads to human review with the evidence collected so far.

The configured terminal cost of referral is 12 units. For assessment, this cost or the loss of an automatic report is added to the diagnostic costs already incurred. Human review is an explicit terminal outcome; the model does not simulate the reviewer's eventual diagnosis.

An audit record associates each action with its outcome, selection score, belief update, cumulative cost, and stopping decision. In the current prototype, diagnostic feedback is simulated using a perturbed likelihood matrix. Moreover, the 40 evaluation cases are sampled from the same mapped pool used to estimate the priors and stage likelihoods. Hiding their labels during an episode therefore does not constitute an independent holdout evaluation. Posterior confidence remains conditional on the assumed model, and the 90% threshold does not establish 90% empirical accuracy.

4 Probability Model and Decision Rule

The probability model represents the agent's uncertainty about a CI failure, while the decision rule determines what to do with that uncertainty. The agent maintains a distribution over possible failure categories and updates it as diagnostic evidence becomes available. This use of a belief state follows the treatment of partially observable decision problems [Kaelbling *et al.*, 1998]. In this project, the hidden category remains fixed during an investigation; diagnostic actions gather information without repairing the underlying fault.

4.1 Prior belief and initial evidence

Let $s^\star$ denote the unknown category, with $\mathcal{S} = \{\text{Code}, \text{Test}, \text{Dependency}, \text{CI/Config}\}$. The four categories

follow the project's mapping of historical labels. They provide one operational label per incident rather than a model of simultaneous faults. Let E_t contain the normalized failed stage z and the t action–outcome pairs observed so far. The belief is

$$b_t(s) = P(s^\star = s \mid E_t), \qquad b_t(s) \geq 0, \qquad \sum_{s \in \mathcal{S}} b_t(s) = 1. \tag{19}$$

The historical records originate from the replication package of Zheng et al. [Zheng *et al.*, 2026]. The project retains 309 of the 375 records, with category counts $(54, 120, 91, 44)$ in the order given above. The remaining 66 records are outside the selected mapping. Let n_s be a category count and $n_{s,z}$ its count at stage z. The smoothed prior and stage likelihood are

$$\widehat{P}(s) = \frac{n_s + \alpha}{N + \alpha|\mathcal{S}|}, \tag{20}$$

$$\widehat{P}(z \mid s) = \frac{n_{s,z} + \beta}{n_s + \beta|\mathcal{Z}|}, \tag{21}$$

where $N = 309$, $\alpha = \beta = 1$, and $|\mathcal{Z}| = 7$. The implemented stage vocabulary includes Build, Test, Setup / Dependency, Quality / Analysis, Deploy / Publish, Ambiguous, and Unknown. The prior vector is therefore $(55, 121, 92, 45)/313$. These probabilities describe the mapped source records, not the prevalence of failures across all CI environments.

The failed stage supplies the initial evidence:

$$b_0(s) = \frac{\widehat{P}(z \mid s)\widehat{P}(s)}{\sum_{u \in \mathcal{S}} \widehat{P}(z \mid u)\widehat{P}(u)}. \tag{22}$$

This separates the prior frequency of a category from the evidence supplied by the current failure. The stage is incorporated once at initialization.

4.2 Observation model and posterior updates

Each diagnostic action a returns a binary outcome $o \in \{0, 1\}$, corresponding to `SIGNAL_NOT_FOUND` and `SIGNAL_FOUND`, respectively. Define the action likelihoods as

$$\ell_a(1 \mid s) = \theta_{s,a}, \qquad \ell_a(0 \mid s) = 1 - \theta_{s,a}. \tag{23}$$

Here $\theta_{s,a}$ is the probability of finding the action-specific signal when the true category is s. These parameters are declared assumptions, because the historical workbook does not contain the diagnostic outcomes needed to estimate them. A signal is evidence about a category, rather than an observation of the category itself.

The model assumes that a new outcome depends on the hidden category and selected action, with no additional dependence on previously observed evidence. Under this assumption, the predictive outcome probability and posterior are

$$q_t(o \mid a) = \sum_{s \in \mathcal{S}} b_t(s)\ell_a(o \mid s), \tag{24}$$

$$b_t^{a,o}(s) = \frac{b_t(s)\ell_a(o \mid s)}{q_t(o \mid a)}. \tag{25}$$

The posterior is evaluated for outcomes with positive predictive probability. After observing o_t from the executed action a_t, the agent sets $b_{t+1} = b_t^{a_t,o_t}$. Both positive and negative results can change the diagnosis. For example, an absent dependency signal reduces belief in a dependency failure when that absence is unlikely under the dependency hypothesis. Correlated signals or repeated use of the same evidence can produce excessive confidence under the independence assumption.

4.3 Expected reporting loss

The most probable category is not necessarily the preferred report when different errors carry different consequences. Following a decision-theoretic approach to troubleshooting [Heckerman *et al.*, 1995], let $L(s, d)$ be the loss of reporting category d when the true category is s. The configured loss can be written compactly as

$$L(s, d) = w_s \, \mathbf{1}\{s \neq d\}, \qquad \boldsymbol{w} = (40, 25, 45, 65). \tag{26}$$

The weights follow the category order given above. Each weight includes the category-specific miss cost and the fixed wrong-report add-on of 5. A correct report has zero terminal loss. The weights are project assumptions expressed in relative loss units.

The expected loss of report d and the minimum reporting risk are

$$r(d \mid b) = \sum_{s \in \mathcal{S}} b(s)L(s, d), \tag{27}$$

$$R(b) = \min_{d \in \mathcal{S}} r(d \mid b). \tag{28}$$

For this loss structure, the minimum-risk report has a useful closed form:

$$\begin{aligned} d_t^{\text{risk}} &\in \arg\max_{d \in \mathcal{S}} w_d b_t(d), \\ R(b_t) &= \sum_{s \in \mathcal{S}} w_s b_t(s) - \max_{d \in \mathcal{S}} w_d b_t(d). \end{aligned} \tag{29}$$

Thus a category with a larger miss cost can become the preferred report before it becomes the most probable category. Ties between reporting categories follow the fixed category order. Human review is excluded from $R(b)$, which measures the loss of the best root-cause report only.

4.4 Selecting evidence by expected risk reduction

Before executing an action, the agent considers both possible outcomes and the report it would prefer after each one. This is an expected-utility calculation over hypothetical observations, as used in Bayesian experimental design [Rainforth *et al.*, 2024]. For an action with positive diagnostic cost $c(a)$, define

$$\overline{R}_t(a) = \sum_{o \in \{0,1\}} q_t(o \mid a)R(b_t^{a,o}), \tag{30}$$

$$\Delta R_t(a) = R(b_t) - \overline{R}_t(a), \tag{31}$$

$$V_t(a) = \frac{\Delta R_t(a)}{c(a)}. \tag{32}$$

Zero-probability outcomes contribute zero to the expectation. The reporting decision is optimized separately within each outcome branch. Under the specified probability model, $\Delta R_t(a)$ is nonnegative because additional evidence can be ignored when choosing a report. However, an individual observed outcome can increase reporting risk even when the expected reduction is positive.

The score measures expected reduction in reporting loss per unit of diagnostic effort. The EIG-per-cost comparison instead measures expected entropy reduction per unit cost. The proposed score is a one-step heuristic: it does not optimize complete action sequences or directly compare every test with the cost of immediate human review. A positive score also does not establish a positive net benefit after subtracting test cost on a common utility scale.

4.5 Reporting, further testing, and human review

Define $d_t^{\mathrm{MAP}} \in \arg\max_s b_t(s)$. The reporting indicator is

$$G_t = \mathbf{1}\{d_t^{\mathrm{MAP}} = d_t^{\mathrm{risk}},\ b_t(d_t^{\mathrm{MAP}}) \geq 0.90,\ R(b_t) \leq 8\}. \tag{33}$$

The three requirements are conjunctive. The risk cap is measured in loss units, not percent. With the configured weights, 90% confidence already implies MAP–minimum-risk agreement and report risk no greater than $65(1 - 0.90) = 6.5$. The agreement and risk-cap checks therefore provide no additional filtering at these settings, although the asymmetric loss continues to influence action selection at lower confidence.

Let $\mathcal{U}_t$ contain unused actions and B_t be the remaining diagnostic budget. After t tests, the feasible and positive-score sets are

$$\mathcal{F}_t = \{a \in \mathcal{U}_t : a \text{ is eligible for } z,\ c(a) \leq B_t,\ t < 8\}, \tag{34}$$

$$\mathcal{F}_t^+ = \{a \in \mathcal{F}_t : V_t(a) > 0\}, \tag{35}$$

where $B_0 = 11$. When $\mathcal{F}_t^+$ is nonempty, select $a_t^\star \in \arg\max_{a \in \mathcal{F}_t^+} V_t(a)$, breaking ties by larger $\Delta R_t(a)$, lower cost, and then action name. For the decision state $x_t = (b_t, z, \mathcal{U}_t, B_t, t)$, the architecture specifies

$$\pi(x_t) = \begin{cases} \mathrm{REPORT}(d_t^{\mathrm{risk}}), & G_t = 1, \\ \mathrm{TEST}(a_t^\star), & G_t = 0 \text{ and } \mathcal{F}_t^+ \neq \varnothing, \\ \mathrm{HUMAN_REVIEW}, & \text{otherwise.} \end{cases} \tag{36}$$

Following a test, the posterior is updated, its cost is deducted from B_t, and the action is removed from $\mathcal{U}_t$. The reporting condition is checked again, including after the eighth permitted test. The positive-score restriction in this rule follows the diagram; the inspected notebook ranks feasible actions but does not yet enforce that restriction. Its existing evaluation therefore does not establish the behavior of this additional stopping condition.

For an investigation ending after T tests, the recorded total cost is

$$C_{\mathrm{episode}} = \sum_{t=0}^{T-1} c(a_t) + \begin{cases} L(s^\star, d_T), & \text{if a category is reported,} \\ 12, & \text{if the case is referred for human review.} \end{cases} \tag{37}$$

The true category is needed only to score a completed automatic report; it is not revealed during the diagnostic episode. Nevertheless, the prototype estimates its prior and stage likelihoods from the same mapped pool from which its 40 evaluation cases are sampled, so those cases do not constitute an independent holdout. Posterior confidence is conditional on the model and its assumed likelihoods; a 90% reporting threshold does not by itself demonstrate 90% empirical accuracy.

5 Test Method

The experiment tests decision behavior under controlled model mismatch rather than replaying raw logs or executing commands against live pipelines. Table 3 summarizes the complete configuration.

5.1 Cases, outcomes, and paired comparison

The 40 labels are hidden from the policy but remain in the same 309-record pool used to estimate priors and stage likelihoods. The evaluation is therefore in-sample with hidden episode labels, not repository-disjoint validation. Since the source records do not contain action outcomes, the harness perturbs the agent likelihoods and pre-generates a shared outcome bank:

$$\widetilde{\theta}_{s,a} = \mathrm{clip}_{[0.02,0.98]}(0.82\theta_{s,a} + 0.09 + \varepsilon_{s,a}), \tag{38}$$

$$\varepsilon_{s,a} \sim \mathcal{N}(0, 0.045^2), \qquad O_{i,a} \sim \mathrm{Bernoulli}(\widetilde{\theta}_{y_i,a}). \tag{39}$$

Every policy receives the same outcome for the same case–action pair. This paired design isolates policy choices within the fixed run, but it does not replace repeated simulations or observed investigation traces.

The three sequential policies use identical actions, likelihoods, gates, and budgets. Risk/cost ranks actions by Equation (32); EIG/cost ranks them by expected entropy reduction per cost; lowest cost chooses the cheapest eligible test; and stage-MAP immediately reports the largest component of b_0. The current implementation does not enforce the formal positive-score guard.

5.2 Measures and checks

Let A_i^π indicate an automatic report and $K_i^\pi = \mathbf{1}\{\widehat{y}_i^\pi = y_i\}$. For $n = 40$,

$$\mathrm{Cov}_\pi = \frac{1}{n}\sum_i A_i^\pi, \tag{40}$$

$$\mathrm{SelAcc}_\pi = \frac{\sum_i A_i^\pi K_i^\pi}{\sum_i A_i^\pi}, \tag{41}$$

$$\mathrm{OCA}_\pi = \frac{1}{n}\sum_i A_i^\pi K_i^\pi. \tag{42}$$

Selective accuracy is reported with coverage because abstention can improve conditional accuracy by accepting fewer cases [Geifman and El-Yaniv, 2017]. Review rate is $1 - \mathrm{Cov}_\pi$. Macro precision, recall, and F_1 are computed over the four states, with referrals counted as false negatives for their true state.

The per-case total cost is

$$T_i^\pi = \sum_{a \in \mathcal{T}_i^\pi} c(a) + \mathbf{1}_{\mathrm{report}} L(y_i, \widehat{y}_i^\pi) + 12\,\mathbf{1}_{\mathrm{review}}. \tag{43}$$

Table 3: Configuration of the deterministic evaluation.

Item	Setting
Source	375 GitHub Actions failures from 260 projects [Zheng *et al.*, 2026]
Mapped pool	309 records: 54 Code, 120 Test, 91 Dependency, 44 CI/Config
Cases	40 balanced cases (10 per state), sampled without replacement
Visible input	Case identifier and normalized failed stage; label hidden during the episode
Policies	Risk reduction/cost, EIG/cost, lowest cost, and stage-MAP
Report gate	MAP equals minimum-risk label, confidence ≥ 0.90, risk ≤ 8
Resources	At most 8 actions, budget 11, review cost 12
Outcome seeds	Likelihood perturbation: 20260826; outcome bank: 20260827

Probability quality is measured with the multiclass Brier score [Brier, 1950] and ten-bin ECE [Guo *et al.*, 2017]:

$$\mathrm{BS}_\pi = \frac{1}{n}\sum_i \sum_{s\in\mathcal{S}} \left(p^\pi_{is} - \mathbf{1}\{y_i = s\}\right)^2, \tag{44}$$

$$\mathrm{ECE}_\pi = \sum_{m=1}^{10} \frac{|\mathcal{B}_m|}{n} \left|\mathrm{acc}(\mathcal{B}_m) - \mathrm{conf}(\mathcal{B}_m)\right|. \tag{45}$$

With only 40 cases, ECE is descriptive rather than evidence that the reporting threshold is calibrated. Fresh-kernel checks verified normalized probabilities, non-repeated actions, stage filtering, common outcome reuse, budget accounting, and terminal report-or-review behavior. The run generated 160 terminal records and 564 action records.

6 Results

Table 1 reports the central safety–coverage trade-off. Risk reduction/cost automatically reported 18 cases: 17 correctly and one incorrectly. Its 45.0% coverage therefore corresponded to 94.44% selective accuracy. EIG/cost reached the same coverage but made two errors. Lowest cost made no wrong automatic report, but referred 25 of 40 cases. Stage-MAP covered every case while making 22 wrong reports.

The cost and probability measures do not identify one universally best policy as mentioned in Table 2. Risk reduction/cost used 4.325 actions and 6.50 diagnostic-cost units per case. Its total decision cost was 589.0, which was 8.4% below EIG/cost, 2.4% below lowest cost, and 42.3% below stage-MAP. It also had the largest macro-F_1 (0.4958) and smallest ECE (0.0387). Lowest cost had the smallest Brier score (0.2357) and perfect selective accuracy, but only 37.5% coverage. These cost values are conditional on the declared error, action, and referral costs; they are not time or money.

Coverage was uneven across states. The proposed policy correctly reported 2 of 10 Code, 0 of 10 Test, 9 of 10 Dependency, and 6 of 10 CI/Config cases. It referred the remaining 8 Code, 10 Test, 1 Dependency, and 3 CI/Config cases, and mislabeled one CI/Config case as Dependency. Thus, the aggregate selective accuracy reflects strong Dependency and CI/Config performance combined with substantial referral, not uniformly reliable diagnosis.

All unit and invariance checks passed in a fresh execution. This supports internal reproducibility of the saved run, but not generalization: one balanced sample, one outcome bank, and the same mapped pool for initialization and sampling provide no confidence interval or independent repository test.

7 Failure Analysis

Incorrect reports and referrals have different meanings. A wrong automatic report exposes the user to a false diagnosis; referral avoids forced guessing but consumes human capacity. Across the sequential policies, three automatic decisions were wrong, corresponding to two unique cases as mentioned in Table 4. Lowest cost had no wrong report at 37.5% coverage, whereas stage-MAP bypassed the gate and made 22 errors.

Table 1: Terminal outcomes on the 40 balanced cases. Selective accuracy is conditioned on automatic reports; OCA is the correct-automatic fraction of all cases.

Policy	Auto.	Correct	Wrong	Review	Coverage (%)	Sel. acc. (%)	OCA (%)
Risk reduction/cost	18	17	1	22	45.0	94.44	42.5
EIG/cost	18	16	2	22	45.0	88.89	40.0
Lowest cost	15	15	0	25	37.5	100.00	37.5
Stage-MAP	40	18	22	0	100.0	45.00	45.0

Table 2: Effort, configured total decision cost, and probability quality.

Policy	Mean actions	Mean diag. cost	Total cost	Macro F_1	Brier	ECE
Risk reduction/cost	4.325	6.5000	589.0	0.4958	0.2673	0.0387
EIG/cost	4.575	6.8500	643.0	0.4703	0.3024	0.0863
Lowest cost	5.200	7.5875	603.5	0.4722	0.2357	0.0668
Stage-MAP	0.000	0.0000	1020.0	0.3824	0.6682	0.1222

Table 4: All incorrect automatic reports made by the sequential policies.

Policy	Case	Stage	True → report	Confidence	Total cost
Risk/cost	CI-007	Quality/Analysis	CI/Config → Dependency	0.9334	68.5
EIG/cost	CI-007	Quality/Analysis	CI/Config → Dependency	0.9334	68.5
EIG/cost	CI-012	Build	Code → CI/Config	0.9052	48.5

7.1 High-confidence errors

The proposed policy's only error, `CI-007`, followed two negative signals. For a negative outcome from action a, posterior odds change by

$$\frac{b_{t+1}(s_i)/b_{t+1}(s_j)}{b_t(s_i)/b_t(s_j)} = \frac{1-\theta_{s_i,a}}{1-\theta_{s_j,a}}. \tag{46}$$

The missing CI/configuration signal had likelihood 0.831 under Dependency but only 0.028 under the true CI/Config state, a Bayes factor of 29.68 in favor of Dependency. The next negative code/test-change signal removed most remaining Code and Test mass. As Figure 2 shows, the resulting Dependency posterior reached 93.34% and passed both gates. The error therefore reflects likelihood misspecification, not a failure to apply Bayes' rule.

`CI-012`, observed only under EIG/cost, contained conflicting signals. A positive code/test-change result initially favored Code; later positive CI/environment evidence and two negative checks shifted the posterior to 90.52% CI/Config. A fixed 0.90 gate cannot prevent error when the observation model itself is overconfident. Calibration must be assessed on independent data [Guo *et al.*, 2017], particularly because probability quality can degrade under distribution shift [Ovadia *et al.*, 2019].

7.2 Referral pattern and model weaknesses

The proposed policy referred 55% of cases, including every Test case and eight Code cases. Referral therefore prevented forced reports but created a large, state-dependent human workload. It should be evaluated as a joint human–model decision rather than counted as an ordinary classification error [Mozannar and Sontag, 2020].

Because Bayesian experimental design is conditional on its probability and utility models [Rainforth *et al.*, 2024], changing the likelihoods, loss matrix, review cost, or threshold can change both action order and policy ranking. External repository splits, real action outcomes, repeated paired runs, dependence-aware evidence models, and observed reviewer performance are required before an operational claim.

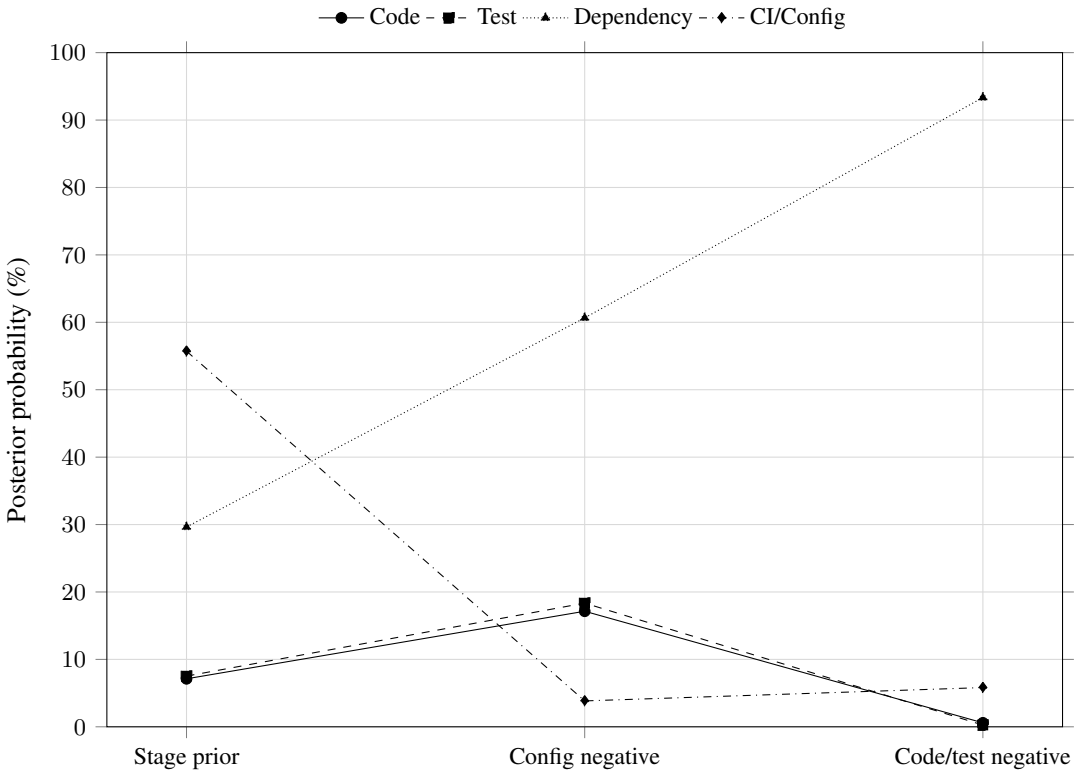


Figure 2: Posterior trajectory for `CI-007`.

8 Conclusion

This study shows how CI diagnosis can be represented as an inspectable sequence of evidence-selection, belief-update, reporting, and referral decisions. In the configured 40-case simulation, risk reduction per cost achieved 94.44% selective accuracy at 45.0% coverage and the lowest total decision cost, while the stage-only baseline made 22 wrong reports. The result supports selective evidence gathering over forced stage-based guessing inside this simulator.

The conclusion remains limited by modeled binary outcomes, declared costs and likelihoods, in-sample initialization, one fixed outcome bank, and uneven coverage. In particular, all Test cases were referred and one wrong report crossed the gate at 0.9334 confidence. The next evaluation should use observed diagnostic traces and repository-disjoint cases, assess calibration and risk–coverage with uncertainty intervals, and measure the receiving human review process. Until then, the agent should remain read-only and advisory, with people retaining authority over diagnosis and remediation.

References

[Brandt *et al.*, 2020] Carolin E. Brandt, Annibale Panichella, Andy Zaidman, and Moritz Beller. LogChunks: A data set for build log analysis. In *Proceedings of the 17th International Conference on Mining Software Repositories*, MSR '20, pages 583–587. Association for Computing Machinery, 2020.

[Brier, 1950] Glenn W. Brier. Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1):1–3, 1950.

[Geifman and El-Yaniv, 2017] Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems*, volume 30, pages 4878–4887. Curran Associates, Inc., 2017.

[Guo *et al.*, 2017] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1321–1330. PMLR, 2017.

[Hassan *et al.*, 2023] Foyzul Hassan, Na Meng, and Xiaoyin Wang. UniLoc: Unified fault localization of continuous integration failures. *ACM Transactions on Software Engineering and Methodology*, 32(6):1–31, 2023. Article 136.

[Heckerman *et al.*, 1995] David Heckerman, John S. Breese, and Koos Rommelse. Decision-theoretic troubleshooting. *Communications of the ACM*, 38(3):49–57, 1995.

[Huang and Lin, 2026] Jiangnan Huang and Bin Lin. On the reruns of GitHub Actions workflows. *ACM Transactions on Software Engineering and Methodology*, 2026. Advance online publication.

[Huang *et al.*, 2026] Yonghui Andie Huang, Daniel Alencar da Costa, Grant Dick, Mariam El Mezouar, and Liwen Xiao. Is this build failure related to my patch? an empirical study of unrelated build failures in continuous integration. *Empirical Software Engineering*, 31(6), 2026. Article 148.

[Kaelbling *et al.*, 1998] Leslie Pack Kaelbling, Michael L. Littman, and Anthony R. Cassandra. Planning and acting in partially observable stochastic domains. *Artificial Intelligence*, 101(1–2):99–134, 1998.

[Lindley, 1956] D. V. Lindley. On a measure of the information provided by an experiment. *The Annals of Mathematical Statistics*, 27(4):986–1005, 1956.

[Mozannar and Sontag, 2020] Hussein Mozannar and David Sontag. Consistent estimators for learning to defer to an expert. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 7076–7087. PMLR, 2020.

[Ovadia *et al.*, 2019] Yaniv Ovadia, Emily Fertig, Jie Ren, Zachary Nado, D. Sculley, Sebastian Nowozin, Joshua V. Dillon, Balaji Lakshminarayanan, and Jasper Snoek. Can you trust your model's uncertainty? evaluating predictive uncertainty under dataset shift. In *Advances in Neural Information Processing Systems*, volume 32, pages 13991–14002. Curran Associates, Inc., 2019.

[Rainforth *et al.*, 2024] Tom Rainforth, Adam Foster, Desi R. Ivanova, and Freddie Bickford Smith. Modern Bayesian experimental design. *Statistical Science*, 39(1):100–114, 2024.

[Shahin *et al.*, 2017] Mojtaba Shahin, Muhammad Ali Babar, and Liming Zhu. Continuous integration, delivery and deployment: A systematic review on approaches, tools, challenges and practices. *IEEE Access*, 5:3909–3943, 2017.

[Valenzuela-Toledo *et al.*, 2025] Pablo Valenzuela-Toledo, Chuyue Wu, Sandro Hernandez, Alexander Boll, Roman Machacek, Sebastiano Panichella, and Timo Kehrer. Explaining GitHub Actions failures with large language models: Challenges, insights, and limitations. In *2025 IEEE/ACM 33rd International Conference on Program Comprehension (ICPC)*, pages 286–297. IEEE, 2025.

[Xu *et al.*, 2025] Weiyuan Xu, Juntao Luo, Tao Huang, Kaixin Sui, Jie Geng, Qijun Ma, Isami Akasaka, Xiaoxue Shi, Jing Tang, and Peng Cai. LogSage: An LLM-based framework for CI/CD failure detection and remediation with industrial validation. In *2025 40th IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 3742–3753. IEEE, 2025.

[Zhang *et al.*, 2022] Chen Zhang, Bihuan Chen, Xin Peng, and Wenyun Zhao. BuildSheriff: Change-aware test failure triage for continuous integration builds. In *Proceedings of the 44th International Conference on Software Engineering*, ICSE '22, pages 312–324. Association for Computing Machinery, 2022.

[Zheng *et al.*, 2026] Lianyu Zheng, Shuang Li, Xi Huang, Jiangnan Huang, Bin Lin, Jinfu Chen, and Jifeng Xuan. Why do GitHub Actions workflows fail? an empirical study. *ACM Transactions on Software Engineering and Methodology*, 35(5):1–29, 2026. Article 139.

[Ziftci *et al.*, 2026] Celal Ziftci, Ray Liu, Spencer Greene, and Livio Dalloro. LLM-based automated diagnosis of integration test failures at Google. In *Proceedings of the IEEE/ACM 48th International Conference on Software Engineering: Software Engineering in Practice*, ICSE-SEIP '26, pages 174–185. Association for Computing Machinery, 2026.